

A1:- Time 23:-

Process	P1	P3	P5	P7	P8
state.	Read from disk unit 3. (R3)	R2	Exits (time expires)	W3	Suspended Not given

which it tries to follow

Let Read from disk unit $i = R_i$
or Write to disk unit $i = W_i$

So, at time 23, P1 is in ready state & reading from disk unit 3.
P3 is also in ready state & it is reading from disk unit 2.
P5 has exited from mem its time expired, so it is suspended.
P7 is in waiting state as it is waiting for P1 to leave resource R3 so that it can write to disk unit 3.
P8 is suspended.
→ Here P7 is waiting or in blocked state and is waiting for relinquish of disk unit 3 by P1.

Time 37:-

P1 read is completed, so its suspended.
P3 → ready state & reading from disk unit 3. read is completed, so suspended.
P5 → Swapped out, so exited.
P7 → ready state & writing to disk unit 3.
P8 → still suspended.

Time 47:-

P1 → suspended

P3 → suspended.

P5 → As P5 was swapped, its interrupt at time 40 was for invalid transaction, so it will be put in queue for writing to disk unit 3 which is still held by P7, so it stays blocked.

P7 → ready state & writing to disk unit 3

P8 → Exited

Matrix:-

	P1	P3	P5	P7	P8
22	Ready	Ready	Suspended	Blocked (P1 resource)	Suspended
37	Suspended	Suspended	Exited	Ready	Suspended
47	Suspended	Suspended	Blocked (P7 resource)	Ready	Exited

106

A2 - \rightarrow Optimal is the best page-replacement algorithm that we can have. So, it gets '1'. ✓

LRU (~~inefficient~~ in implementation note)

✓ \rightarrow LRU is the second best algorithm although we don't implement it due to constraints (efficiency).

✓ \rightarrow 2nd-chance replacement is an approximate algorithm for LRU which is implemented using FIFO stack so it's the next best option.

✓ \rightarrow FIFO is the worst among the given options.

Considering a scale of 1 for perfect & 5 for bad, we can say that none of them are close to optimal in real-world scenario & the rating is

Optimal \rightarrow 1

LRU \rightarrow 2

2nd-chance \rightarrow 3

FIFO \rightarrow 5

10

FIFO performs worst in these & it also suffers from Belady's anomaly among these replacement algorithms.

How about rest of the algorithms ??

Q3:- string : 1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6.

LRU → replaces least recently used page.

FIFO → ^{use} replaces the first in first out principle

Optimal → Uses future view to replace the page which will occur the farthest in future.

Algo Frame size	1	2	3	4	5	6	7
1) LRU	20 ✓	18 ✓	15 ✓	10 ✓	8 ✓	7 ✓	7 ✓
2) FIFO	20 ✓	18 ✓	16 ✓	14 ✓	16 ✓	8 ✓	7 ✓
3) Optimal	20 ✓	15 ✓	11 ✓	8 ✓	7 ✓	7 ✓	7 ✓

(12)

Ans:- Given semaphore variables S & R, so a semaphore variable is acquired using semWait & released using semSignal.

```
foo() {  
  do {  
    1 → semWait(S);  
    2 → semWait(R);  
    3 → x++;  
    4 → semSignal(S);  
    5 → semSignal(R);  
    3 while(1);  
  }  
}
```

```
bar() {  
  do {  
    1 → semWait(R);  
    2 → semWait(S);  
    x--;  
    semSignal(S);  
    semSignal(R);  
  } while(1);  
}
```

a.) Consider line 1 of both functions. So, foo() acquires S & bar() acquires R, if such a thing happens during the concurrent execution of foo() & bar(), then neither foo() will be able to complete its line 2 nor bar() will be able to complete its line 2. foo() wants R which is held by bar() & bar() wants S which is held by foo(). So, a deadlock occurs here & both processes are being blocked forever until one of them is forced to relinquish the ~~its~~ semaphore. (12)

b.) No, concurrent execution of these 2 processes can't result in indefinite postponement of one of them. Definite postponement of either will happen. e.g. Consider foo() has reached line 3. but bar() is at line 1 (not able to complete the instruction, but foo() will keep going and as soon as 4 & 5 line of foo() is completed, bar() will acquire R and S and so bar() will be executed & foo() will wait. Indefinite postponement of either func will not occur, although both together can be blocked forever as in part (a).

26, 37, 100, 14, 88, 33, 99, 12.

AS:- a.) FCFS :-

$$\text{Total movement} = (37-26) + (100-37) + (100-14) + (88-14) + (88-33) + (99-33) + (99-12)$$

$$= 11 + 63 + 86 + 74 + 55 + 66 + 87$$

$$= 442.$$

b.) SSTF :-

$$\text{Total movement} = \{ 26 \rightarrow 33 \rightarrow 37 \rightarrow 14 \rightarrow 12 \rightarrow 88 \rightarrow 99 \rightarrow 100$$

$$= 7 + 4 + 23 + 2 + 76 + 11 + 1$$

$$= 124.$$

c.) SCAN (going up) :- (Assuming cylinder length of 1 to 100).

$$\text{Total movement} = 26 \rightarrow 33 \rightarrow 37 \rightarrow 88 \rightarrow 99 \rightarrow 100 \rightarrow 14 \rightarrow 12$$

$$= 7 + 4 + 51 + 11 + 1 + 86 + 2$$

$$= 162.$$

d.) C-SCAN (going up) :- (Assume cylinder length of 1 to 100). Total length = 100.

$$\text{Total movement} = 26 \rightarrow 33 \rightarrow 37 \rightarrow 88 \rightarrow 99 \rightarrow 100 \rightarrow 1 \rightarrow 12 \rightarrow 14$$

(went to beginning)

$$= 7 + 4 + 51 + 11 + 1 + 99 + 11 + 2$$

$$= 186.$$

Note, if we don't consider movement from one end to other, i.e., from 100 to 1 cylinder, the movement of 99, then total movement = $186 - 99 = 87$.

(Considering going up means increasing cylinder value, i.e., from 26 \rightarrow 100)

11

Ans:- Rotation speed = 15000 rpm

$$\begin{aligned}\text{Avg. rotation latency} &= \frac{1}{2} \left(\frac{60}{15000} \right) \text{ s} \\ &= \frac{1}{2} \times \frac{1}{250} \times 1000 \text{ ms} = 2 \text{ ms}.\end{aligned}$$

$$\begin{aligned}\text{b.) Avg. access time} &= \text{Avg. seek time} + \text{avg rotation latency} \\ &= 4 \text{ ms} + 2 \text{ ms} = 6 \text{ ms}\end{aligned}$$

now, Tracks = 1000
Sectors / track = 400

$$\begin{aligned}\text{bytes / sector} &= 512 \\ \text{File size} &= 1 \text{ MByte} \\ &= (1024)^2 \text{ Byte} \\ &= 2^{20} \text{ Bytes} = 2^2 \times 8 \text{ bytes} \\ &= 2^{23} \text{ bytes}.\end{aligned}$$

a.) Total transfer time :-

$$\begin{aligned}\text{Total bytes in a track} &= 512 \times 400 \text{ bytes} \\ &= 2^9 \times 2^2 \times 100 \text{ bytes} \\ &= 2^{11} \times 100 \text{ bytes} = 25 \times 2^{13} \text{ bytes}.\end{aligned}$$

$$\text{Total tracks required} = \frac{2^{23}}{25 \times 2^{13}} = \frac{2^{10}}{25} = 40.96$$

\therefore Total of 41 tracks will be required for contiguous allocation on disk out of which 40 will be filled completely & 41th track will be partially filled.

Total transfer time = Avg. seek time to access 41 tracks
+ rotation time to completely access 40 tracks
+ rotation time for partial access to 41 tracks

$$\begin{aligned}&= \cancel{4 \text{ ms}} \times 41 + 2 \text{ ms} \times 40 + 2 \text{ ms} \times 0.96 \\ &= 164 \text{ ms} + 160 \text{ ms} + 1.92 \text{ ms} \\ &= 325.92 \text{ ms}\end{aligned}$$

[as 41th track only 0.96 of it needs to be written].

(Assuming avg. seek time denotes avg. time to seek from one track position to another).

c.) Rotational delay is already calculated = 2ms.

d.) Total time to read 1 sector = Avg. seek time + line ^{to a track} to read a sector on a track.

$$= 4ms + \frac{1}{400} \times 2ms = 4ms + \frac{1}{200}ms = 4.005ms$$

e.) Total time to read 1 track = Avg. access time + Avg. rotation delay

$$= 4ms + 2ms = 6ms.$$

Ans:-

Since a direct pointer can hold a block of 8KB, an indirect pointer can hold a block of $(8 \times 8)KB = 64KB$, double indirect can hold $(8 \times 8 \times 8)KB = 512KB$ & triple indirect pointer can hold $(8 \times 8 \times 8 \times 8)KB = 4096KB$.

\therefore Total (max) file size the given i-node can hold

$$= 13 \times 8KB + 64KB + 512KB + 4096KB$$

$$= (104 + 64 + 512 + 4096)KB$$

$$= 4776KB.$$

(6)

why 16??

$$\text{Total size} = 13 \times 8KB + 1024 \times 16 \times 8KB + (1024 \times 16)^2 \times 8KB + (1024 \times 16)^3 \times 8KB$$

\therefore 32-bit pointer \rightarrow 8KB = $\frac{8 \times 8 \times 8 \times 1024}{32} = 1024 \times 16$

$$\frac{8KB}{32b} = \frac{8 \times 1024 \times 8}{32} = 256 \times 8 = 2048 = \underline{\underline{2048}}$$

A8:- 1) The target can keep ID of ^{all the} ~~incoming~~ ^{challenges} ~~connections~~ ^{connections} so that are being given to others or is getting from others. So, there can be many ^{challenges} ~~connections~~ corresponding to different connections with the target. If any challenge is repeated (of asked challenges and being challenges asked to target), break the connection with that. ~~connect~~

16

Now, it may happen that the target has got same ~~same~~ challenge from another ~~other~~ connection which he has asked to attacker. In such a case, we don't want to break the connection. So, we will let the target answer the challenge of valid / invalid connection wrong whenever same challenge appears from other side.

If we get same answer as reply to our same challenges, then its attacker & so the connection be broken, else if we get correct response, we create connection.

→ Another approach is to keep ~~in~~ track of number of connections between any pair, ~~if~~ and simply reject more than one connection b/w a pair.

This won't work

→ OR ~~OR~~ Allow more than one connection b/w a pair and keep track of questions asked, if question is same, then reject the connection. The target will be determined in such a case by looking at timestamp and who asked the same question first.

END